

A note on Endiansim

Foundation of High Performance Computing @ UniTS, 2020-2021

Intro

You know that “a basic-type variable” is nothing else but a bunch of n bytes in memory whose bits are read *as if* they form the given basic type (i.e. `int`, `float`, `double`, ...).

However, there is one more detail to be considered when that memory region spans multiple bytes: is the most significant byte stored either in the lowest- or in the highest- addresses in memory ?

To clarify the meaning of that let’s consider how *we*, the humans, write numbers in the commonly used arabic notation considering the large number 3,141,592,653.

Since we write from left to right, the lowest “memory” (i.e. the position in the page starting from the top-left corner which has coordinate 0) position is occupied by the first digit `3` which is *the most significant one*:

| | | | | | | | | | | | |
|------------------|---|---|---|---|---|---|---|---|---|---|----|
| “Memory” address | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
| | 3 | 1 | 4 | 1 | 5 | 9 | 2 | 6 | 5 | 3 | 9 |

Hence, in the arabic notation the most significant digits occupy the lowest memory addressed and as the significance (i.e. the positional value, or in other words the power of 10 to which the digit corresponds) decreases the “address in the page” increases.

In other words, the *Big-end* of the number we want to write is stored at the begin of the memory area we used to store it: this notation is referred to as *big-endianism*.

At the opposite, in the case we wrote from right to left (still considering the top-left corner as the origin) we would have that the *least significant* digits would occupy the lowest memory addresses:

| | | | | | | | | | | | |
|------------------|---|---|---|---|---|---|---|---|---|---|----|
| “Memory” address | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
| | 9 | 3 | 5 | 6 | 2 | 9 | 5 | 1 | 4 | 1 | 3 |

That notation, in which the most significant digits are stored at the end of its memory region, is called *little-endianism*.

So, given that both ways are equally legitimate, how are the number represented in a computer? That is an architectural choice and it amounts to an ab-initio CPU’s architectural choice.

It happens that both ways are present in the world: some CPU implement little-endianism (like Intel’s and AMD’s *x86* architecture) and some other implement big-endianism (like IBM’s *powerPC* architectures and many historical RISC architectures like the old Motorola 68k). More exotic treatment are also present (ARM’s are generally little-endian but, depending on the version, allow switching). In both little- and big- endianism the *bit-order* inside a single byte is always little-endian (the most significant bit is the the highest address in the byte).

There existed also examples of the so called *middle-endianism* that were big-endian systems which also swapped to big-endian bits-order in each byte.

The big-endian choice is understandable in the light of the frequent necessity in the past to look at memory dumps by human sight, events in which a human-like big-endian order was also a clear advantage (I would dare to say that real *digital natives* are those who understand that by experience; in any case, remember, when you to look ad a little-endian memory dump, that you must reverse the order of byte interpretation).

As a side consideration, the human big-endian power-of-ten representation that we use as western humans may have been motivated by the necessity of getting immediately the most significant part of a number, which in our left-to-right convention comes to be lies at the lowest “addresses”.

Practical consideration

Let's then examine the code that I gave to you to handle this matter, in the `swpa_image()` routine.

For instance, `0x100` is `1 00000000` in binary and hence it sets the first bit of the second byte and no bits in the first byte; `0xFF` is `11111111` in binary and hence it sets all the bits in the first byte. The logical `and` between the two results in the two following sequences.

[illegible]

Little-endian

0x100

| lowest byte | | | | | | | | | highest byte | | | | | | | | |
|-------------|---|---|---|---|---|---|---|---|--------------|---|---|---|---|---|---|---|---|
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

0xFF

| lowest byte | | | | | | | | | highest byte | | | | | | | | |
|-------------|---|---|---|---|---|---|---|---|--------------|---|---|---|---|---|---|---|---|
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 |
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

0x100 & 0xFF

| lowest byte | | | | | | | | | highest byte | | | | | | | | |
|-------------|---|---|---|---|---|---|---|---|--------------|---|---|---|---|---|---|---|---|
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

And so, clearly, checking the result of the `&` you can easily understand the endianness of the architecture you're running on. In the actual code that I gave to you, the value `0xF` is used. It sets the first 4 bits of the lowest byte, but since we are checking for zero/non-zero condition that is perfectly equivalent to the case depicted above.

How to adapt your bytes to an alien endianness

Switching between big- and little- endianness is fairly simple. It amounts to swap bytes. When the bytes are only two, as in the `pgm` image format, it is the simplest case.

Using the bit-shift and bit-wise logical operators the task is accomplished (`MEM` in the following is the `short int` value we want to swap):

1. `MEM & (short int)0xff00` builds a `short int` value that has only the most significant byte of the original `MEM`. The most significant byte is still in its original position.
2. `(MEM & (short int)0xff00) >> 8` builds the value and shifts it by 1 byte *to the right*, i.e. the shift moves the most significant byte in the least significant position. It is the same than dividing the original value by 256; actually that first operation could be simplified in `(MEM >> 8)`.
3. `(MEM & (short int)0x00FF) << 8` selects the least significant byte and shifts it to the most significant position. Also this operation could be written as `(MEM << 8)`.
4. summing the 2 previous results, which could also be written as a bit-wise or `|`, gives you back your swapped 2-bytes values.

The simplifications `MEM >> 8` and `MEM << 8` have not been adopted for educational purposes. They can be used only in the 2-bytes case, while if you have 4, 8 or even more bytes you have to select individually each byte by using suited bit masks.

Are you willing to write your own general routine to swap endianness?

At the end of the course I could provide you my own for testing purposes.